



# Core Energy Comparison

Out-of-Order



In-Order



Dally, IEEE Computer'08

Do we always need the extra hardware?

# MorphCore: An Energy-Efficient Architecture for High-Performance ILP and High-Throughput TLP

Khubaib\*

M. Aater Suleman<sup>\*+</sup>

Milad Hashemi<sup>\*</sup>

Chris Wilkerson<sup>‡</sup>

Yale N. Patt<sup>\*</sup>

<sup>\*</sup> HPS Research Group  
The University of Texas at Austin

<sup>+</sup> Calxeda Inc.

<sup>‡</sup> Intel Labs

Courtesy: Prof. Yale Patt, UT Austin

# The Need for an Adaptive Core

- Sometimes a single thread with high ILP
  - Need a heavy-weight out-of-order core
  - Provides high performance by exploiting ILP
- Sometimes many threads
  - Out-of-order is unnecessary
  - Need a power-efficient core
  - Provides high performance by exploiting thread-level parallelism
- We need an adaptive core that can do both
  - Exploits instruction-level parallelism when needed
  - Exploits thread-level parallelism when needed

# Problem

Current core architectures **do not adapt**

Large cores **limit performance when TLP is high**

Small cores **limit performance when TLP is low**

# Outline

- Problem Statement
- Previous Work
  - Asymmetric chip multiprocessors
  - Reconfigurable core architectures
- MorphCore
- Evaluation

# Asymmetric Chip Multiprocessors

- One or few large out-of-order cores with many small in-order cores

[Morad+ CAL'06, Suleman+ TR'07, Hill+ Computer'07, Suleman+ ASPLOS'09]

- Limited flexibility

- Fixed number of large and small cores

- Migration overhead

- Migrate the thread state/data to large core

# Reconfigurable Core Architectures


- **Fundamental Idea**
  - Build a chip with “simpler cores” and “combine” them at runtime using additional logic to form a high-performance out-of-order core
  - Core Fusion - Ipek+ ISCA'07, TFlex - Kim+ MICRO'07, Federation Cores - Tarjan+ DAC'08, and many others
- **Fused core has low performance and low energy-efficiency**
  - Increased latencies among its pipeline stages
- **Significant mode switching overhead**

# Outline

- Problem Statement
- Previous Work
- MorphCore
  - Key Insights and Basic Idea
  - Design and Operation
- Evaluation



# Key Insight 1: The Potential of In-Order SMT



- With 8 threads, the in-order core's performance almost matches the out-of-order core's

# Key Insight 2

Minimal changes to a traditional OOO core  
can transform it into a  
highly-threaded in-order SMT core

Existing structures in an OOO core  
can be re-used to support  
highly-threaded in-order SMT execution

# MorphCore: Basic Idea

The opposite of previous proposals:

A) The base design: OOO core

B) Then we add in-order SMT

Two modes:

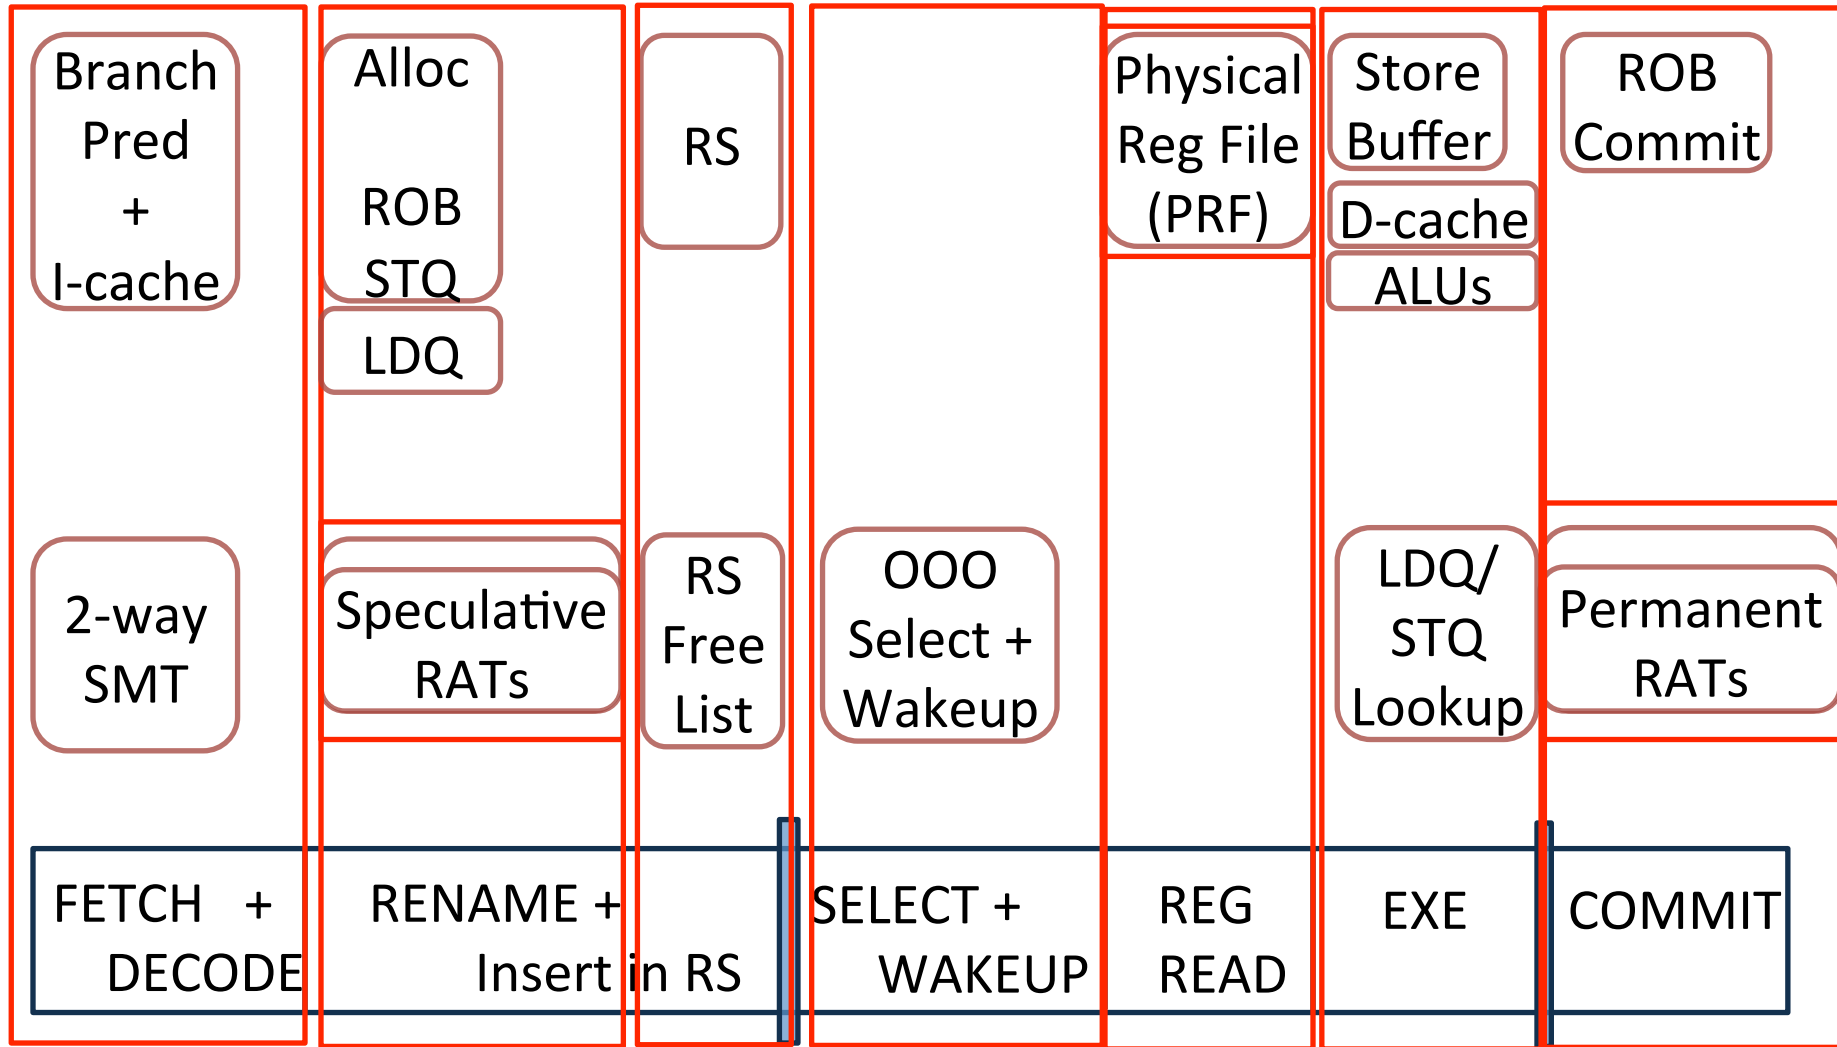
OutOfOrder	out-of-order core Exploits ILP High single-thread performance
------------	---

InOrder	highly-threaded in-order SMT core Exploits TLP High multi-thread performance No OOO execution → Energy savings
---------	---

# Outline

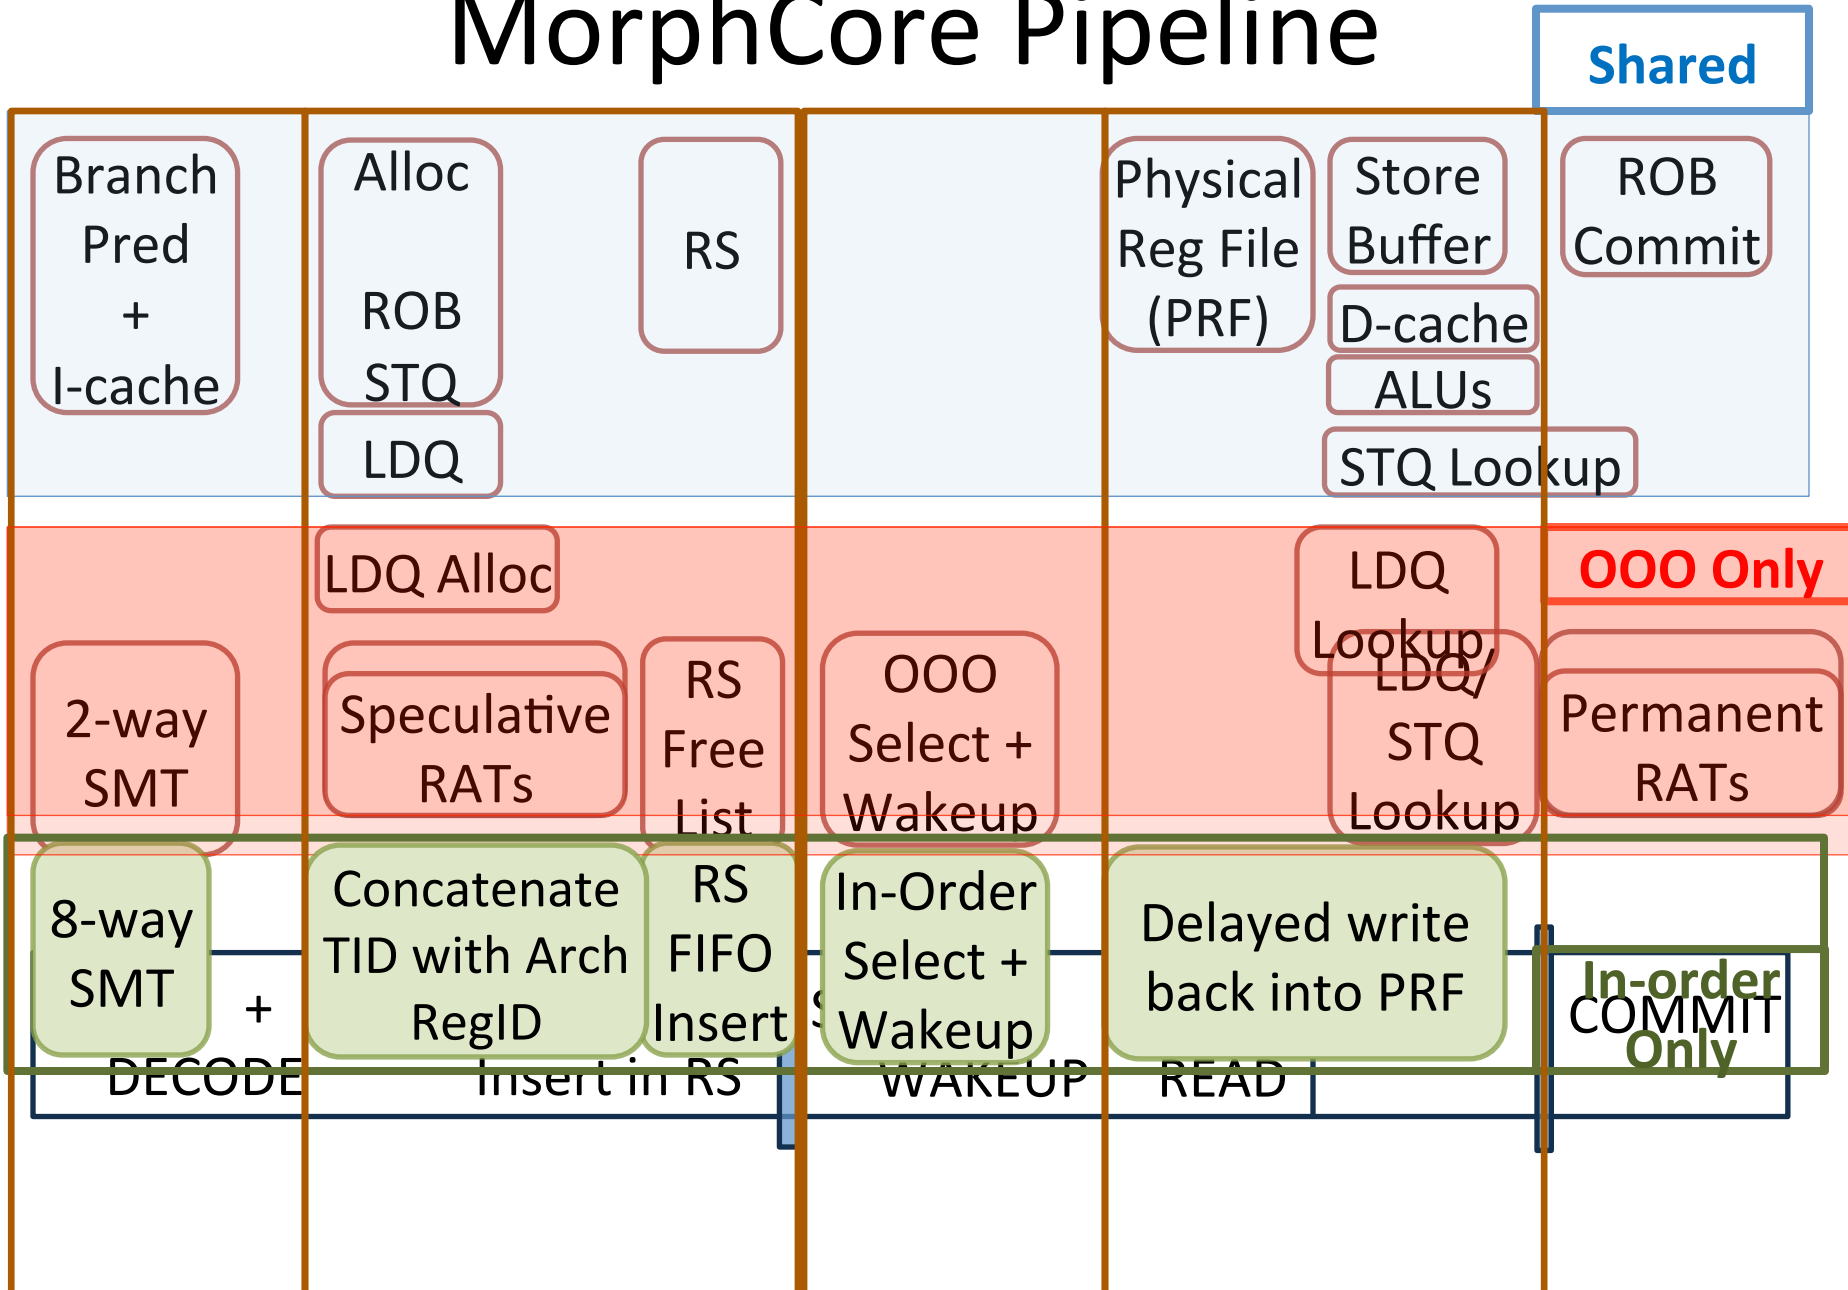
- Problem Statement
- Previous Work
- MorphCore
  - Key Insights and Basic Idea
  - Design and Operation
- Evaluation

# Baseline OOO Pipeline



# MorphCore Pipeline

# MorphCore Pipeline



# Microarchitecture Summary

- Use existing structures without modification
  - Physical Register File (PRF), Decode, Execution pipeline
- Use existing structures with minor modification
  - OOO Reservation Stations → InOrder instruction queues
  - Because of InOrder execution, delayed writeback into PRF (extra bypass)
- SMT related changes
  - Front-end (e.g. multiple PCs, branch history regs), changes in resource allocation algorithms
- In-Order instruction scheduler



# Overheads

- Core area increases by 1.5%
  - Increase in SMT contexts (0.5%)  
(Note that added contexts are in-order, so no additional rename tables and physical registers)
  - InOrder Wakeup and Select Logic (0.5%)
  - Extra bypass (0.5%)
- Core frequency decreases by 2.5%
  - Add multiplexers in the critical path of 2 stages
    - Rename and Scheduling

# Mode Switching Policy

- Number of active threads  $\leq 2$  ?
- OutofOrder when active threads  $\leq 2$ 
  - MorphCore can support up to 2 000 threads
  - TLP is limited so execute OOO to obtain performance
- InOrder when active threads  $> 2$ 
  - More than 2 threads can only run simultaneously in InOrder mode
  - TLP is high so high core throughput and energy savings can be obtained by executing threads in-order

# How Mode Switching Happens?

- (1) **Drains** the core pipeline
- (2) **Spills** architectural registers of **currently active threads** to reserved ways in the private 256KB L2
- (3) **Turns off/on** Renaming, OOO Scheduling, Load Queue
- (4) **Fills** the architectural registers of **next-active threads** into PRF (update RATs when going into OutofOrder)

Currently an overhead of 300 - 450 cycles

# Outline

- Problem Statement
- Previous Work
- MorphCore
- Evaluation

# Methodology

- Detailed cycle-level x86 simulator
- McPAT (modified) to calculate energy/area
- Performance/energy evaluation of MorphCore vs. alternative architectures
  - Large OOO cores: optimized for single-thread
  - Medium and Small cores: optimized for multi-thread
- Workloads
  - Single-threaded (ST): 14 – SPEC 2006
  - Multi-threaded (MT): 14 – Databases, SPLASH, others

# Evaluated Architectures

**All comparisons on *approximately* equal area**

ST : single-thread

MT: multi-thread

OOO : out-of-order

InO : in-order

Core	# of cores	Freq. (GHz)	Type	Issue width	SMT threads Per core	Total threads	Peak throughput ops/cycle	
							ST	MT

# Evaluated Architectures

All comparisons on *approximately* equal area

ST : single-thread

MT: multi-thread

OOO : out-of-order

InO : in-order

Core	# of cores	Freq. (GHz)	Type	Issue width	SMT threads Per core	Total threads	Peak throughput ops/cycle	
							ST	MT
OOO-2	1	3.4	OOO	4	2	2	4	4
OOO-4	1	-5%	OOO	4	4	4	4	4

# Evaluated Architectures

All comparisons on *approximately* equal area

ST : single-thread

MT: multi-thread

OOO : out-of-order

InO : in-order

Core	# of cores	Freq. (GHz)	Type	Issue width	SMT threads Per core	Total threads	Peak throughput ops/cycle	
							ST	MT
OOO-2	1	3.4	OOO	4	2	2	4	4
OOO-4	1	-5%	OOO	4	4	4	4	4
MED	3	3.4	OOO	2	1	3	2	6



# Evaluated Architectures

All comparisons on *approximately* equal area

ST : single-thread

MT: multi-thread

OOO : out-of-order

InO : in-order

Core	# of cores	Freq. (GHz)	Type	Issue width	SMT threads Per core	Total threads	Peak throughput ops/cycle	
							ST	MT
OOO-2	1	3.4	OOO	4	2	2	4	4
OOO-4	1	-5%	OOO	4	4	4	4	4
MED	3	3.4	OOO	2	1	3	2	6
SMALL	3	3.4	InO	2	2	6	2	6

# Evaluated Architectures

All comparisons on *approximately* equal area

ST : single-thread

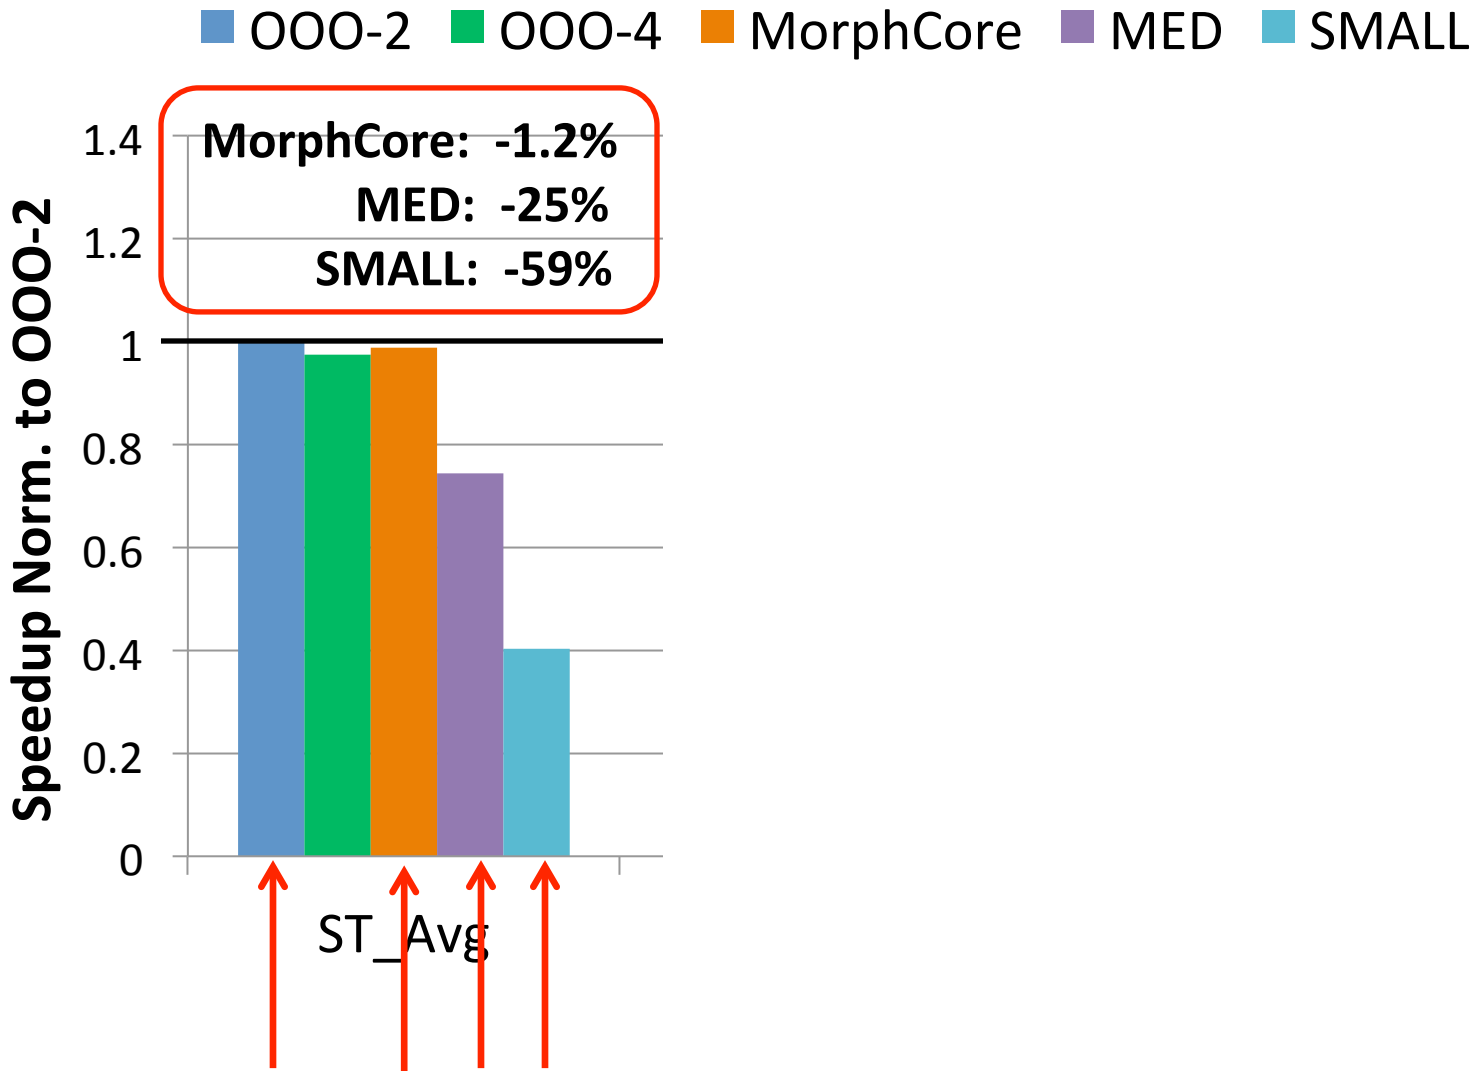
MT: multi-thread

OOO : out-of-order

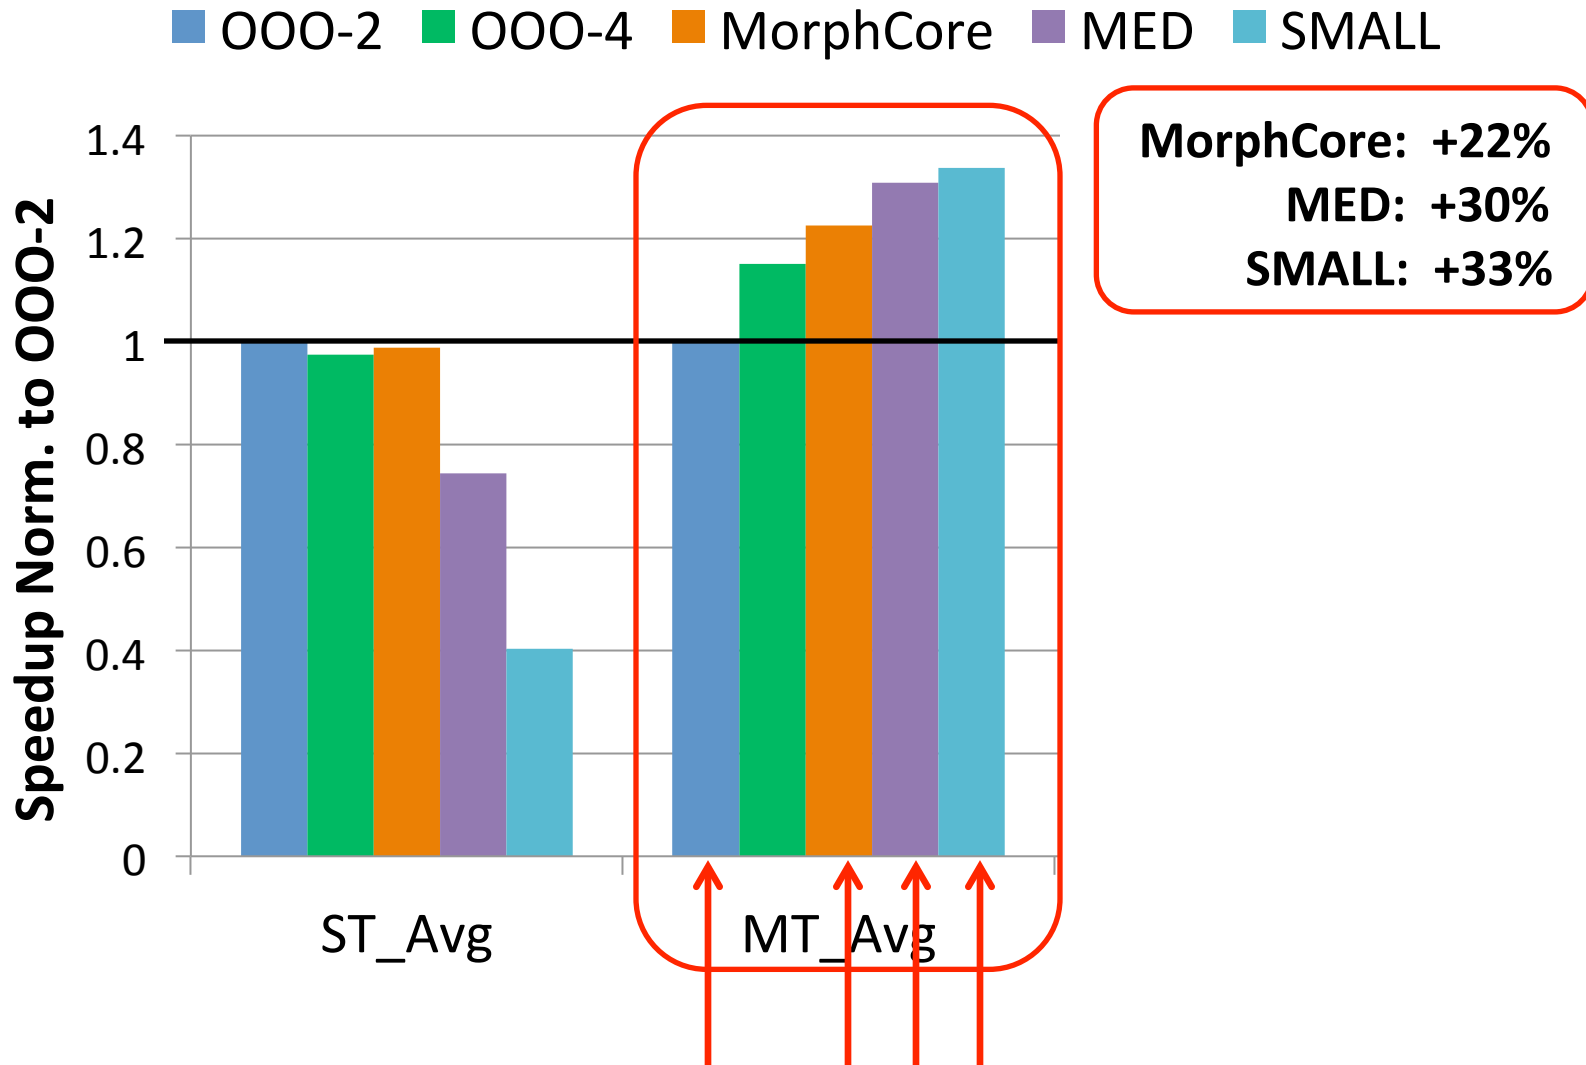
InO : in-order

Core	# of cores	Freq. (GHz)	Type	Issue width	SMT threads Per core	Total threads	Peak throughput (ops/cycle)	
							ST	MT
OOO-2	1	3.4	OOO	4	2	2	4	4
OOO-4	1	-5%	OOO	4	4	4	4	4
MED	3	3.4	OOO	2	1	3	2	6
SMALL	3	3.4	InO	2	2	6	2	6
MorphCore	1	-2.5%	OOO/ InO	4	2 OOO/ 8 InO	2 OOO/ 8 InO	4	4


# Performance: Single-thread




# Performance: Multi-thread



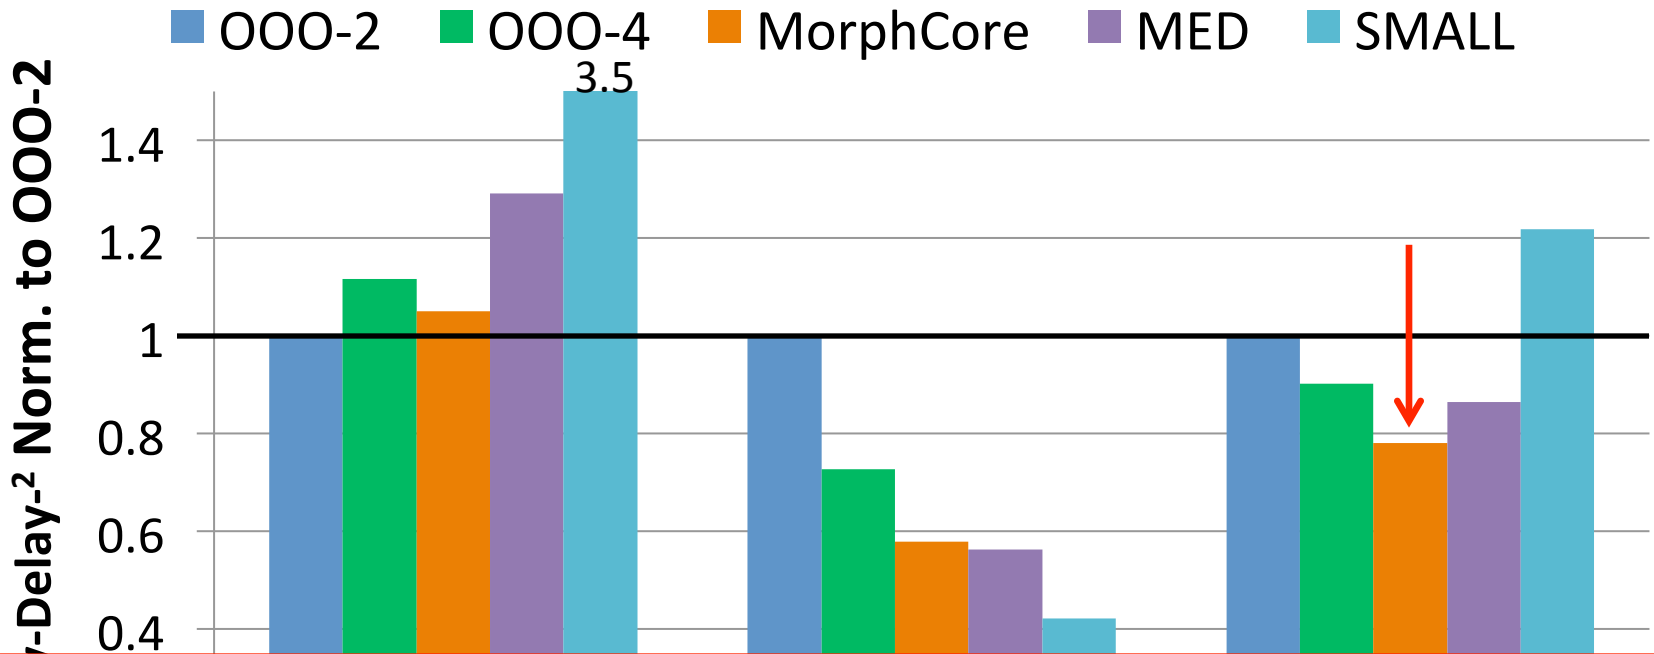
# Performance: Both ST and MT



# Energy



# Energy-delay-squared ( $ED^2$ )



**On average, across all workloads, MorphCore provides the lowest  $ED^2$   
22% lower than OOO-2 and 44% lower than SMALL**

# Summary

- MorphCore **adapts well to both** single-thread and multi-thread workloads
- **Requires minimal changes** to a traditional OOO core
- Operates in two modes:
  - OOO core when TLP is low
  - Highly-threaded in-order SMT core when TLP is high
- Significantly outperforms other alternative architectures